

# TASK-BASED SPARSE CHOLESKY SOLVER ON TOP OF RUNTIME SYSTEM

---

Iain S. Duff, Jonathan D. Hogg and **Florent Lopez**  
Sparse Days, 2016

Rutherford Appleton Laboratory  
**NLAFET Project**

Solve  $Ax = b$ , where  $A$  is **large** and **sparse**, on modern architectures.

Exploiting modern platforms is challenging:

- **Multicore** processors and deep **memory hierarchy**.
- **Heterogeneous** e.g. CPU & GPU or Xeon Phi.
- **Distributed-memory** systems.

Use **Direct Method**: Sparse Cholesky factorization  $A = LL^T$

▲ Numerically robust and general purpose

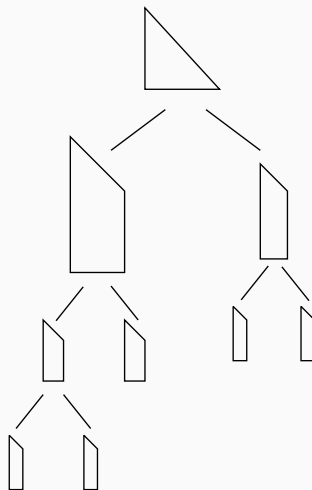
▼ High memory usage and computational cost

# SPARSE CHOLESKY FACTORIZATION

The numerical factorization of  $A$  rely on an *elimination tree* expressing data dependencies in the factor  $L$ . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of  $L$ .

The tree is traversed in a *topological order*, and each node is factorised using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernode.

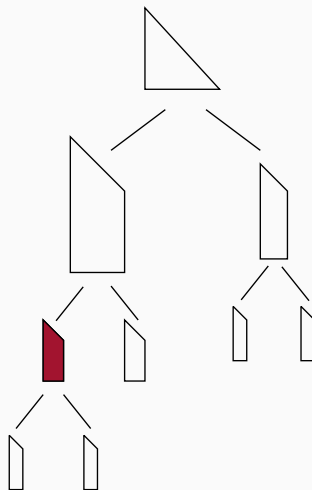


# SPARSE CHOLESKY FACTORIZATION

The numerical factorization of  $A$  rely on an *elimination tree* expressing data dependencies in the factor  $L$ . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of  $L$ .

The tree is traversed in a *topological order*, and each node is factorised using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernode.

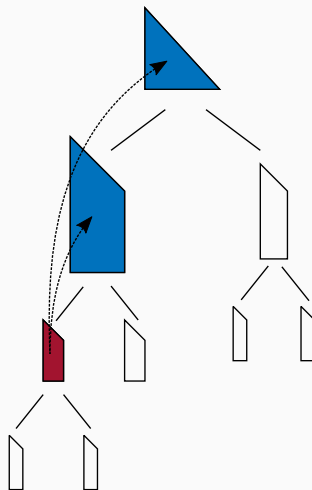


# SPARSE CHOLESKY FACTORIZATION

The numerical factorization of  $A$  rely on an *elimination tree* expressing data dependencies in the factor  $L$ . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of  $L$ .

The tree is traversed in a *topological order*, and each node is factorised using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernode.



Sequential Task Flow (STF) programming model:

- Tasks are submitted to the runtime system following the [sequential algorithm](#).
- The runtime analyses manipulated data and infers task dependencies in order to ensure the [sequential consistency](#) of the parallel code.
- The DAG is executed via a [dynamic scheduling](#) of the (ready) tasks on the architectures.
- The runtime may be capable of automatically handling the data transfer across the architecture.
- [Superscalar analysis](#) in processors: dependency detection between instructions in order to issue them in parallel.

## OpenMP 4.0

- `task` construct and `depend` clause (`in`, `out`, `inout`).
- No control on the scheduling strategy.
- Shared-memory system only.

## StarPU

- `starpup_insert_task`, data *handle* with access mode (`R`, `W`, `RW`).
- Full control on scheduling policy with possibility to implement new one.
- API for distributed-memory systems.

# THE STF SPARSE CHOLESKY FACTORIZATION

```
forall nodes snode in post-order
  call alloc(snode) ! allocate data structures

  call submit(init, snode:W) ! initialize node structure
end do

forall nodes snode in post-order
  ! factorize node
  do k=1..n in snode
    call submit(factorize, snode:R, blk(k,k):RW) ! factorize block

    do i=k+1..m in snode
      call submit(solve, blk(k,k):R, blk(i,k):RW) ! perform solve
    end do

    do j=k+1..n in snode
      do i=k+1..m in snode
        call submit(update, blk(j,k):R, blk(i,k):R, blk(i,j):RW)
      end do
    end do

    forall ancestors(snode) anode
      do j=k+1..p(anode) in snode
        do i=k+1..m in snode
          call submit(update_btw, blk(j,k):R, blk(i,k):R, a_blk(rmap(i), cmap(j)):RW)
        end do
      end do
    end do

  end do
end do
```



# EXPERIMENTS

#	Matrix	Flops ( $10^9$ )	Application/description
1	Schmid/thermal2	18.6	Unstructured thermal FEM
2	Rothberg/gearbox	22.8	Aircraft flap actuator
3	DNVS/m_t1	23.4	Tubular joint
4	DNVS/thread	35.7	Threaded connector
5	DNVS/shipsec1	40.5	Ship section
6	GHS_psdef/crankseg_2	48.8	Linear static analysis
7	AMD/G3_circuit	67.3	Circuit simulation
8	Koutsovasilis/F1	228	AUDI engine crankshaft
9	Oberwolfach/boneS10	297	Bone micro-FEM
10	ND/nd12k	514	3D mesh problem
11	JGD Trefethen/Trefethen_20000	669	Integer matrix
12	ND/nd24k	2080	3D mesh problem
13	Oberwolfach/bone010	3910	Bone micro-FEM
14	GHS_psdef/audikw_1	5840	Automotive crankshaft

- Symmetric positive-definite matrices
- Metis nested dissection ordering
- Machine: 2 x 14 cores E5-2695 v3 (Haswell) @ 2.30GHz

Parametrized Task Graph (PTG) programming model:

- Uses a **compact representation** of the DAG which is problem size independent.
- The dataflow between tasks is **explicitly** encoded (i.e. task dependencies are explicitly given).
- The runtime handles the communications implicitly using the dataflow representation.
- Under some hypothesis, the dataflow information can be **automatically** extracted from the sequential code using a dedicated compiler: **not in our case** unfortunately.