

TASK-BASED SPARSE CHOLESKY SOLVER ON TOP OF RUNTIME SYSTEM

Iain S. Duff, Jonathan D. Hogg and **Florent Lopez**
PMAA'16, 2016

Rutherford Appleton Laboratory
NLAFET Project

Solve $Ax = b$, where A is **large** and **sparse**, on modern architectures.

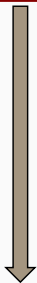
Using **Direct Method**: Sparse Cholesky factorization $A = LL^T$

- ▲ Numerically robust and general purpose
- ▼ High memory usage and computational cost

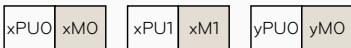
Exploiting modern platforms is challenging:

- **Multicore** processors and deep **memory hierarchy**.
- **Heterogeneous** e.g. CPU & GPU or Xeon Phi.
- **Distributed-memory** systems.


Application



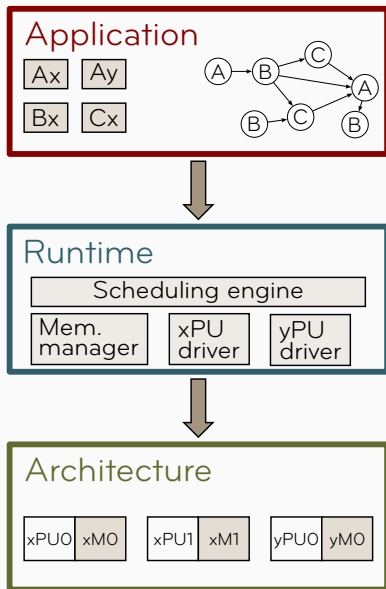
Architecture



- The classical approach is based on a mixture of technologies (e.g., MPI+OpenMP+CUDA) which.
 - programming costs.
 - is difficult to maintain and update.
 - is prone to (performance) portability issues.



- The classical approach is based on a mixture of technologies (e.g., MPI+OpenMP+CUDA) which.
 - programming costs.
 - is difficult to maintain and update.
 - is prone to (performance) portability issues.
- runtimes provide an abstraction layer that hides the architecture details.



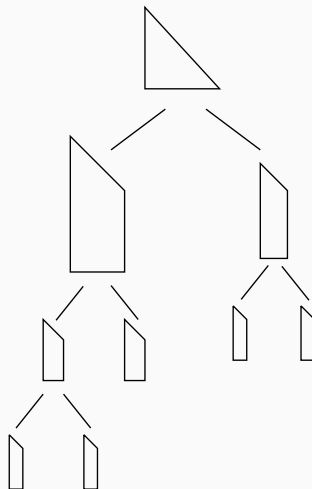
- The classical approach is based on a mixture of technologies (e.g., MPI+OpenMP+CUDA) which.
 - programming costs.
 - is difficult to maintain and update.
 - is prone to (performance) portability issues.
- runtimes provide an abstraction layer that hides the architecture details.
- the workload is expressed as a DAG of tasks.

SPARSE CHOLESKY FACTORIZATION

In numerical factorization of A the *elimination tree* expresses data dependencies in the factor L . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of L .

The tree is traversed in a *topological order*, and each node is factorized using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernodes.

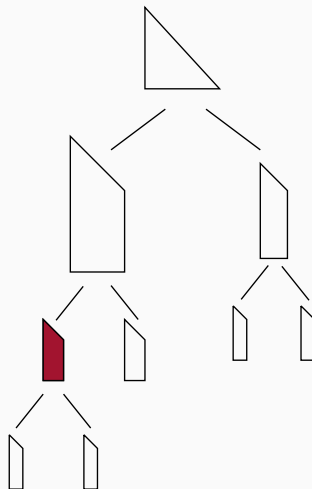


SPARSE CHOLESKY FACTORIZATION

In numerical factorization of A the *elimination tree* expresses data dependencies in the factor L . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of L .

The tree is traversed in a *topological order*, and each node is factorized using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernodes.

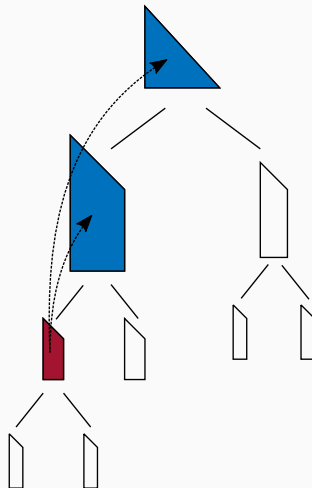


SPARSE CHOLESKY FACTORIZATION

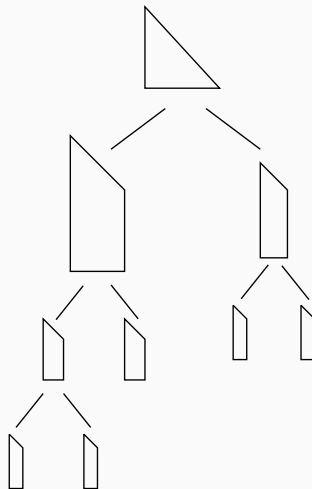
In numerical factorization of A the *elimination tree* expresses data dependencies in the factor L . Each node, referred to as *supernode*, is a *dense* lower trapezoidal *submatrix* of L .

The tree is traversed in a *topological order*, and each node is factorized using *dense Cholesky algorithm*.

Updates between node are handled using a *supernodal scheme* i.e. updates are applied directly to the target supernodes.




Sources of **parallelism** in the
elimination tree:



SPARSE CHOLESKY FACTORIZATION: PARALLELISM

Sources of **parallelism** in the elimination tree:

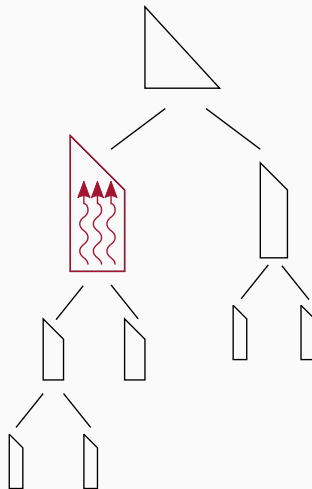
- **Tree parallelism**: Supernode in independent branches can be processed **concurrently**.




SPARSE CHOLESKY FACTORIZATION: PARALLELISM

Sources of **parallelism** in the elimination tree:


- **Tree parallelism**: Supernode in independent branches can be processed **concurrently**.
- **Node parallelism**: When a supernode is large enough, it may be processed in parallel.



TASK-BASED SPARSE CHOLESKY FACTORIZATION

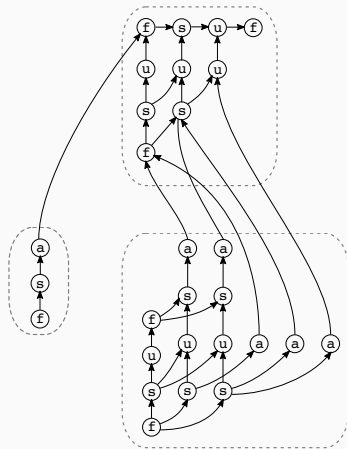
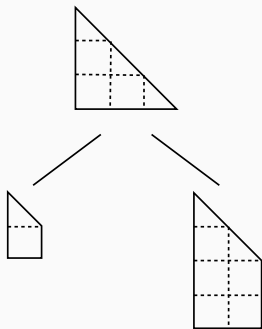


TASK-BASED SPARSE CHOLESKY FACTORIZATION



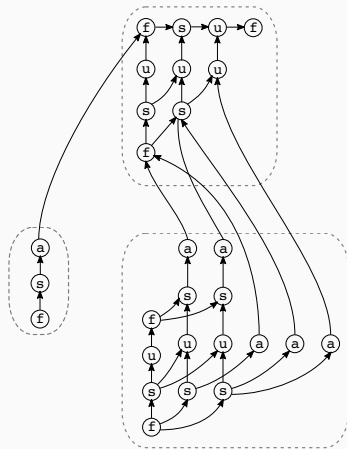

Supernodes are **partitioned** into square **blocks** ($nb \times nb$) on which operations are applied (factorize, solve, update, update_between).

TASK-BASED SPARSE CHOLESKY FACTORIZATION



Supernodes are **partitioned** into square **blocks** ($nb \times nb$) on which operations are applied (**factorize**, **solve**, **update**, **update_between**). The **DAG** replaces the elimination tree for representing the dependencies.

TASK-BASED SPARSE CHOLESKY FACTORIZATION



Supernodes are **partitioned** into square **blocks** ($n_b \times n_b$) on which operations are applied (**factorize**, **solve**, **update**, **update_between**). The **DAG** replaces the elimination tree for representing the dependencies. Implemented in the HSL package [MA87](#).

TASK-BASED SPARSE CHOLESKY FACTORIZATION

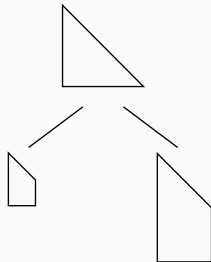
```
forall nodes snode in post-order
  call alloc(snode) ! allocate data structures

  call init(snode) ! initianlize node structure
end do

forall nodes snode in post-order
  ! factorize node
  call factorize(snode)

  ! update ancestor nodes
  forall ancestors(snode) anode
    call update_btw(snode, anode)
  end do

end do
end do
```



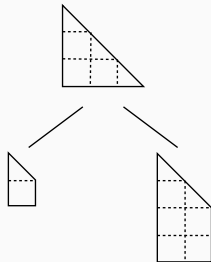
TASK-BASED SPARSE CHOLESKY FACTORIZATION

```
forall nodes snode in post-order
  call alloc(snode) ! allocate data structures

  call init(snode) ! initialize node structure
end do

forall nodes snode in post-order
  ! factorize node
  do k=1..n in snode
    call factorize(blk(k,k)) ! factorize block
    ! solve block
    do i=k+1..m in snode
      call solve(blk(k,k), blk(i,k))
    end do
    ! update block
    do j=k+1..n in snode
      do i=k+1..m in snode
        call update(blk(j,k), blk(i,k), blk(i,j))
      end do
    end do

    ! update ancestor nodes
    forall ancestors(snode) anode
      do j=k+1..p(anode) in snode
        do i=k+1..m in snode
          call update_btw(blk(j,k), blk(i,k),
                        a_blk(rmap(i), cmap(j)))
        end do
      end do
    end do
  end do
end do
```



Sequential Task Flow (STF) programming model:

- In the parallel code, tasks are submitted to the runtime system following the **sequential algorithm**.
- The runtime analyses the manipulated data and infers task dependencies in order to ensure the **sequential consistency** of the parallel code.
- **Superscalar analysis** in processors: dependency detection between instructions in order to issue them in parallel.
- The DAG is executed via a **dynamic scheduling** of the (ready) tasks on the architecture.
- The runtime may be capable of automatically handling the **data transfers** on the architecture (e.g. CPU/GPU memory nodes).