

TASK-BASED SPARSE CHOLESKY SOLVER ON TOP OF RUNTIME SYSTEM

Iain S. Duff, Jonathan D. Hogg and **Florent Lopez**
Sparse Days, 2016

Rutherford Appleton Laboratory
NLAFET Project

Solve $Ax = b$, where A is **large** and **sparse**, on modern architectures.

Exploiting modern platforms is challenging:

- **Multicore** processors and deep **memory hierarchy**.
- **Heterogeneous** e.g. CPU & GPU or Xeon Phi.
- Distributed memory systems.

Use **Direct Method**: Sparse Cholesky factorization $A = LL^T$

▲ Numerically robust and general purpose

▼ High memory usage and computational cost

The numerical factorization of A rely on an *elimination tree* expressing data dependencies in the factor L . Each node, referred to as *supernode*, is a *dense* (lower trapezoidal) *submatrix* of the factor L .

The tree is traversed in a *topological order*, and each node is factorised using *dense algorithm*.

Updates between node is handle using a *supernodal scheme* i.e. updates are applied directly to the target supernode.

Sequential Task Flow (STF) programming model:

- Tasks are submitted to the runtime system following the [sequential algorithm](#).
- The runtime analyses manipulated data and infers task dependencies in order to ensure the [sequential consistency](#) of the parallel code.
- The DAG is executed via a [dynamic scheduling](#) of the (ready) tasks on the architectures.
- The runtime may be capable of automatically handling the data transfer across the architecture.
- [Superscalar analysis](#) in processors: dependency detection between instructions in order to issue them in parallel.

THE STF SPARSE CHOLESKY FACTORIZATION

```
forall nodes snode in post-order
  ! allocate data structures
  call alloc(snode)
  ! initialize node structure
  call submit(init, snode:W)
end do

forall nodes snode in post-order

  ! factorize node
  do k=1..n in snode
    ! factorize diagonal block
    call submit(factorize, snode:R, blk(k,k):RW)

    do i=k+1..m in snode
      ! perform triangular solve w.r.t diag block
      call submit(solve, blk(k,k):R, blk(i,k):RW)
    end do

    do j=k+1..n
      do i=k+1..m
        call submit(update, blk(j,k):R, blk(i,k):R, blk(i,j):RW)
      end do
    end do

    forall ancestors(snode) anode
      ! update ancestor nodes
      call submit(update_between, snode:R, anode:RW)
    end do

  end do
end do
```

Parametrized Task Graph (PTG) programming model:

- Uses a **compact representation** of the DAG which is problem size independent.
- The dataflow between tasks is **explicitly** encoded (i.e. task dependencies are explicitly given).
- Under some hypothesis, the dataflow information can be **automatically** extracted from the sequential code using a dedicated compiler: **not in our case** unfortunately.